

Previous work

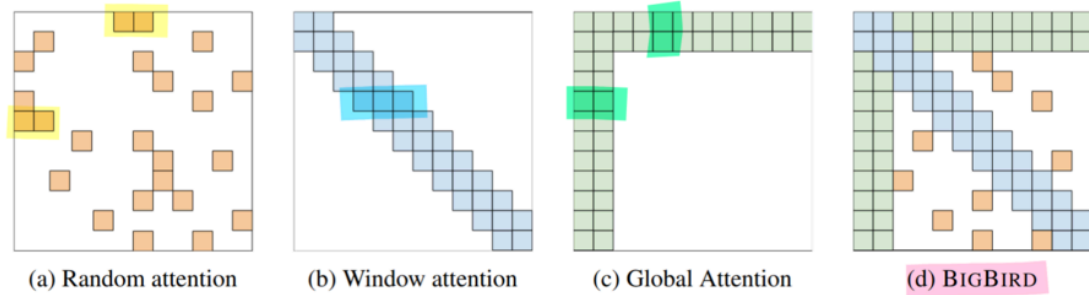


Figure 1: Building blocks of the attention mechanism used in BIGBIRD. White color indicates absence of attention. (a) random attention with $r = 2$, (b) sliding window attention with $w = 3$ (c) global attention with $g = 2$. (d) the combined BIGBIRD model.

Introduction

My goal was to generate python functions, using only math, numpy and pure python, with a algorithm approche, simplicity and easy to port to C. For each attention mask (a, b, c and d), i have written the following features :

- generate a boolean mask, using the corresponding parameter
- same, but based on a given sparsity (by having a function : given sparsity -> corresponding parameter)
- generate artificial matrix of only ones and zero based on the mask
- a test function which shows how to use it and the output

Code

Imports

```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import math # comment
```

UTILS

```
In [2]: def get_nb_non_zero(matrix):
return np.count_nonzero(matrix)
```

```
In [3]: def get_density(matrix, length):
return float(get_nb_non_zero(matrix)) / float(length * length)
```

```
In [4]: def get_sparsity(matrix, length):
return 1.0 - get_density(matrix, length)
```

```
In [5]: def show_matrix_infos(matrix, length, given_sparsity = -1.0): # -1.0 correspond to None.
# conditions : shape(matrix) = (length, length)
```

```

real_sparsity = get_sparsity(matrix, length)
given_text = "None" if given_sparsity < 0 else f"{given_sparsity:.2f}"
text = f"Length: {length}, Given Sparsity: {given_text}, Real Sparsity: {real_sparsity:.4f}"
plt.title(label=text)
plt.imshow(matrix, cmap='gray', interpolation='nearest', vmin=0, vmax=1)
plt.colorbar()
plt.show()

```

RANDOM ATTENTION

By number of non-zeros per row

```

In [6]: def get_random_attention_mask(length, nz_per_row):
# conditions : nz_per_row <= length
rng = np.random.default_rng()
mask = rng.multivariate_hypergeometric([1]*length, nz_per_row, size=length).astype(bool)
return mask

```

```

In [7]: def generate_matrix_with_random_attention_mask(length, nz_per_row):
# conditions : nz_per_row <= length
matrix = np.ones((length, length))
mask = get_random_attention_mask( length=length, nz_per_row=nz_per_row)
matrix[~mask] = 0
return matrix

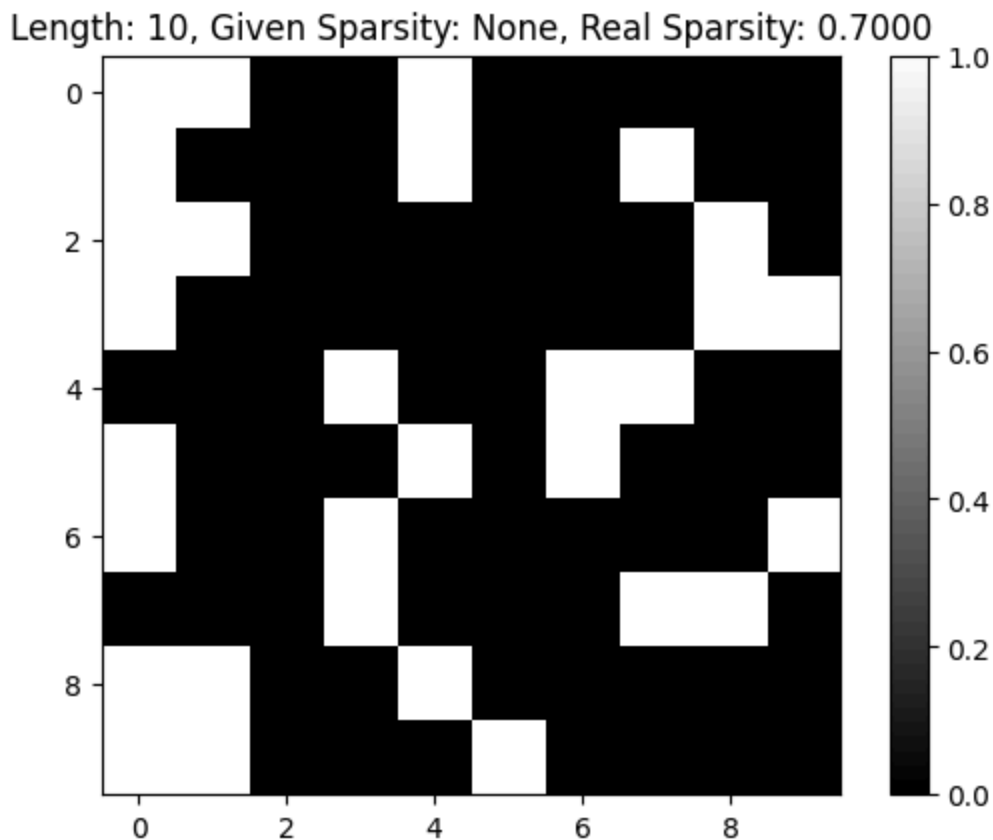
```

```

In [8]: def test_generate_matrix_with_random_attention_mask():
length = 10
nz_per_row = 3
matrix = generate_matrix_with_random_attention_mask(length=length, nz_per_row=nz_per_row)
show_matrix_infos(matrix=matrix, length=length)

test_generate_matrix_with_random_attention_mask()

```



By sparsity

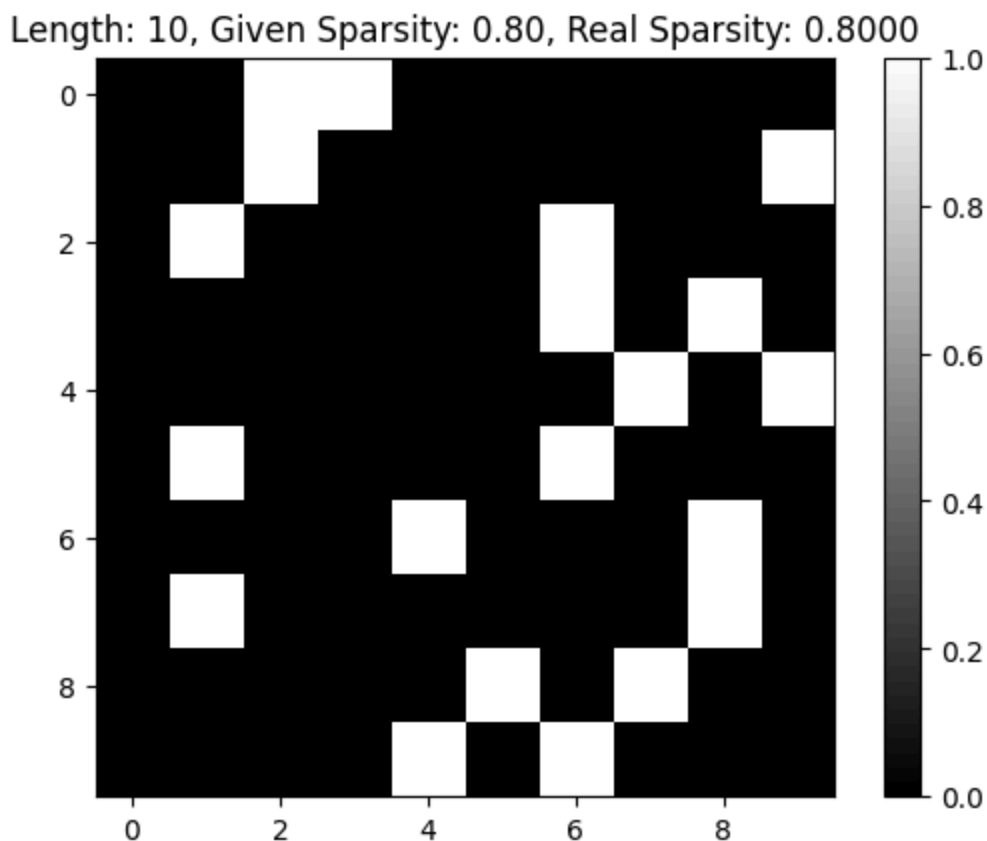
```
In [9]: def best_nz_per_row_from_sparsity(length, sparsity):
# conditions : 0 <= sparsity <= 1
return round(length * (1 - sparsity))
```

```
In [10]: def get_random_attention_mask_with_sparsity(length, sparsity):
# conditions : 0 <= sparsity <= 1
nz_per_row=best_nz_per_row_from_sparsity(length=length, sparsity=sparsity)
return get_random_attention_mask( length=length, nz_per_row=nz_per_row)
```

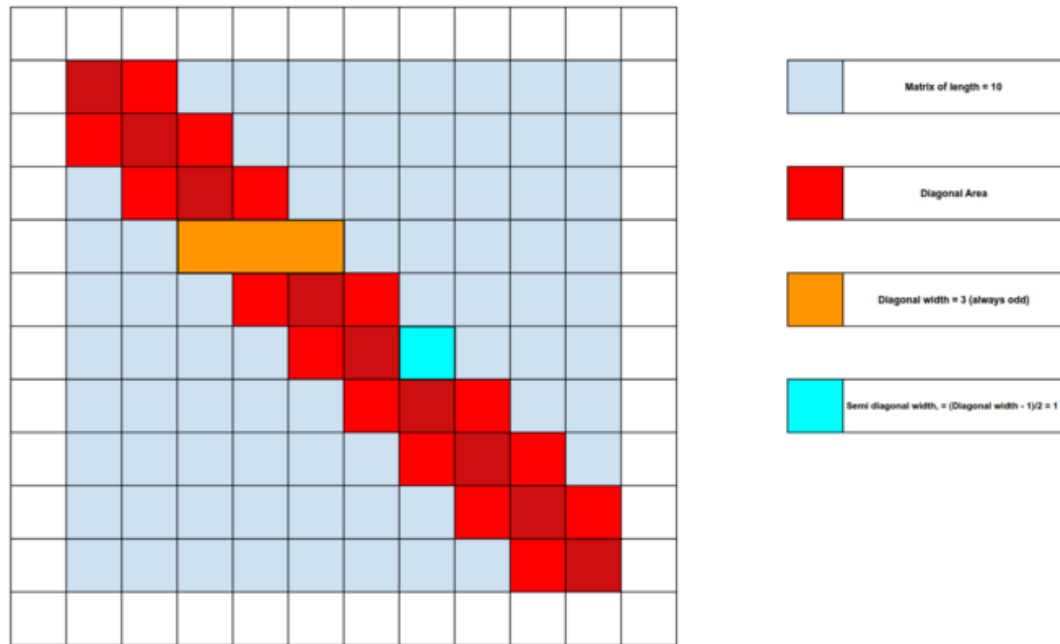
```
In [11]: def generate_matrix_with_random_attention_mask_with_sparsity(length, sparsity):
# conditions : 0 <= sparsity <= 1
matrix = np.ones((length, length))
mask = get_random_attention_mask_with_sparsity(length=length, sparsity=sparsity)
matrix[~mask] = 0
return matrix
```

```
In [12]: def test_generate_matrix_with_random_attention_mask_with_sparsity():
length = 10
sparsity = 0.8
matrix = generate_matrix_with_random_attention_mask_with_sparsity(length=length, sparsity=sparsity)
show_matrix_infos(matrix=matrix, length=length, given_sparsity=sparsity)

test_generate_matrix_with_random_attention_mask_with_sparsity()
```



WINDOW ATTENTION



Utils

```
In [13]: def diagonal_area(length, diagonal_width):
# conditions : length
if diagonal_width == 0:
    return 0
else:
    n = length
    #semi diagonal width
    sdw = diagonal_width // 2 # (diagonal_width / 2 - 1 because is odd)
    da = n * ( 1 + 2 * sdw ) - sdw * (sdw + 1)
    return da
```

By diagonal width

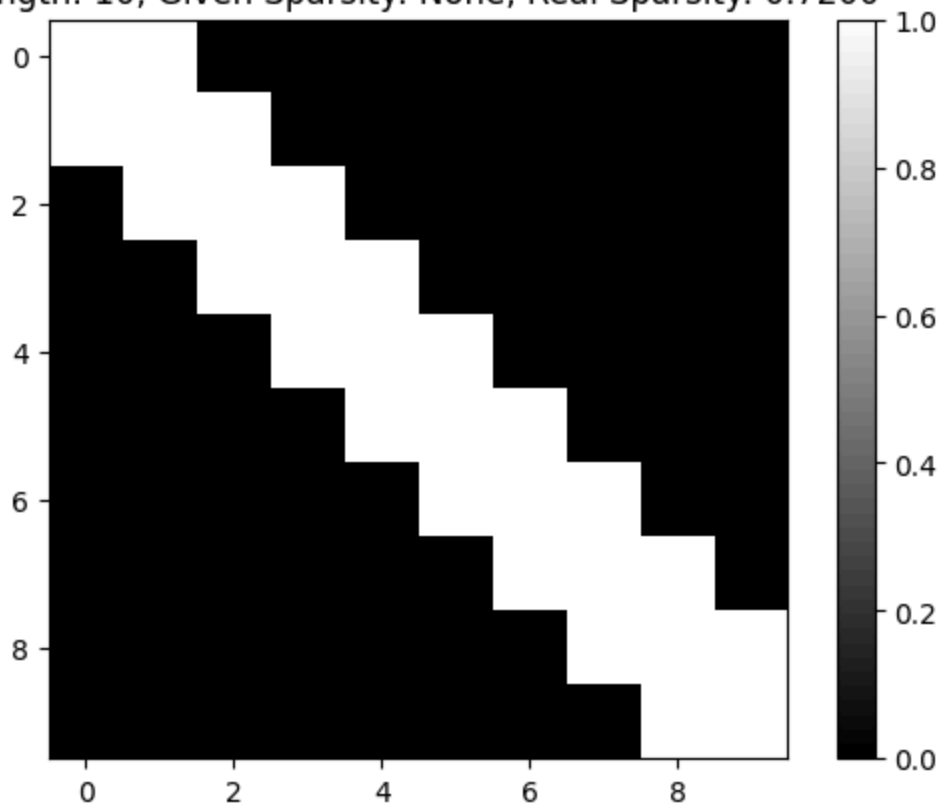
```
In [14]: def get_window_attention_mask (length, diagonal_width):
# conditions : shape(matrix) = (length, length), 0 <= diagonal_width <= 2*length - 1 (cover full matrix), diagonal_width
mask = np.zeros(shape=(length, length), dtype=bool)
if (diagonal_width > 0):
    sdw = diagonal_width // 2
    if diagonal_width == 1:
        mask = np.fromfunction(lambda i, j: j == i, shape=(length, length), dtype=int)
    else :
        mask = np.fromfunction(lambda i, j: np.abs(i - j) <= sdw ,shape=(length, length), dtype=int)
    return mask
```

```
In [15]: def generate_matrix_with_window_attention_mask(length, diagonal_width):
# conditions : 0 <= diagonal_width <= 2*length - 1 (cover full matrix), diagonal_width is odd
matrix = np.ones((length, length))
mask = get_window_attention_mask( length= length, diagonal_width=diagonal_width)
matrix[~mask] = 0
return matrix
```

```
In [16]: def test_generate_matrix_with_window_attention_mask():
length = 10
diagonal_width = 3
matrix = generate_matrix_with_window_attention_mask(length=length, diagonal_width=diagonal_width)
show_matrix_infos(matrix, length)

test_generate_matrix_with_window_attention_mask()
```

Length: 10, Given Sparsity: None, Real Sparsity: 0.7200



By sparsity

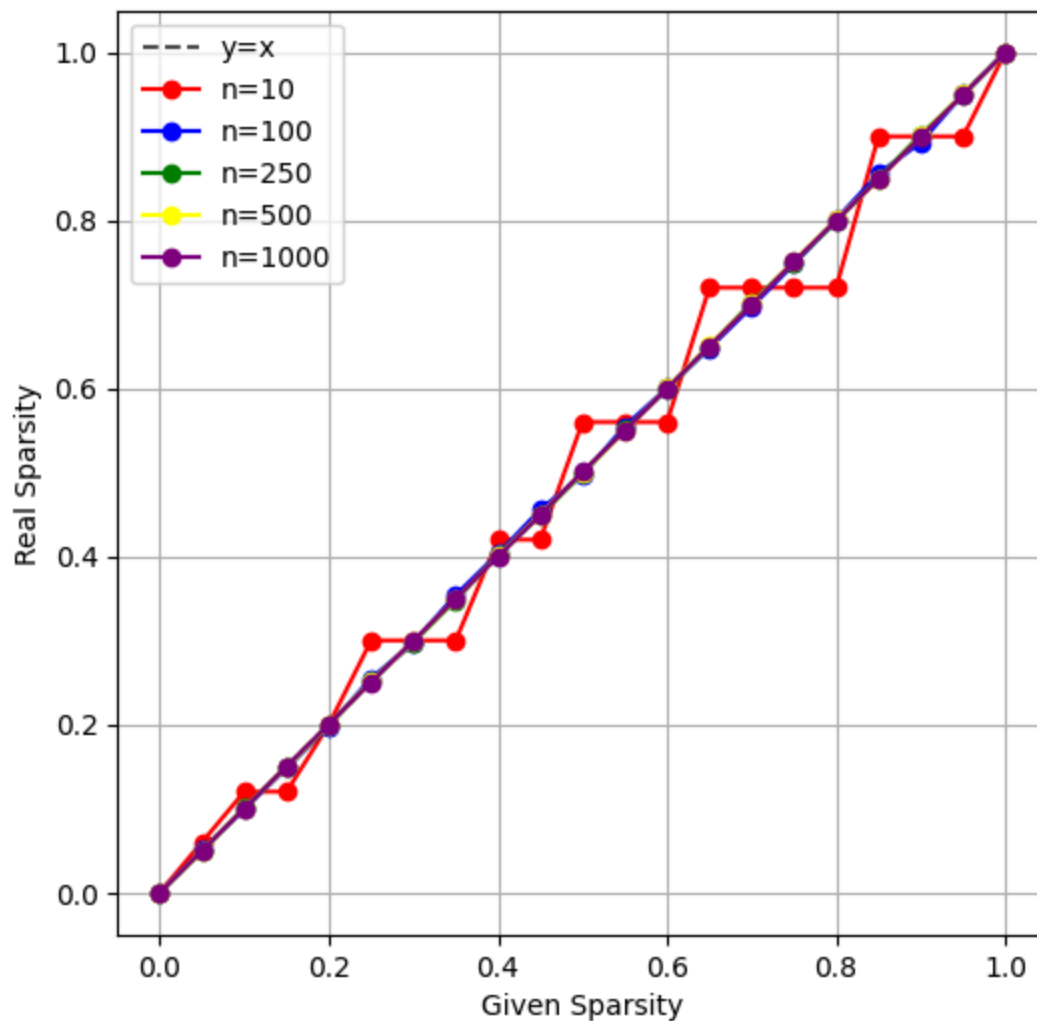
```
In [17]: def best_diagonal_width_from_sparsity(length, sparsity):
n = length
density = 1.0 - sparsity
# ideal diagonal area
da = n * n * density
# from this point, all is explained in the related document
a = -1
b = 2 * n - 1
c = n - da
det = b * b - 4 * a * c
x = (-b + math.sqrt(det)) / (2 * a)

sdw = round(x)

dw = 2 * sdw + 1

if(dw < 0) : dw = 0
elif(dw > 2*n - 1): dw = 2*n - 1
# print(f"For matrix of size: {n} and given sparsity: {sparsity}, ideal semi diagonal width is : {x}, chosen dw is {dw}")
return dw
```

```
In [19]: # test function hided for render
test_diagonal_width_from_sparsity()
```

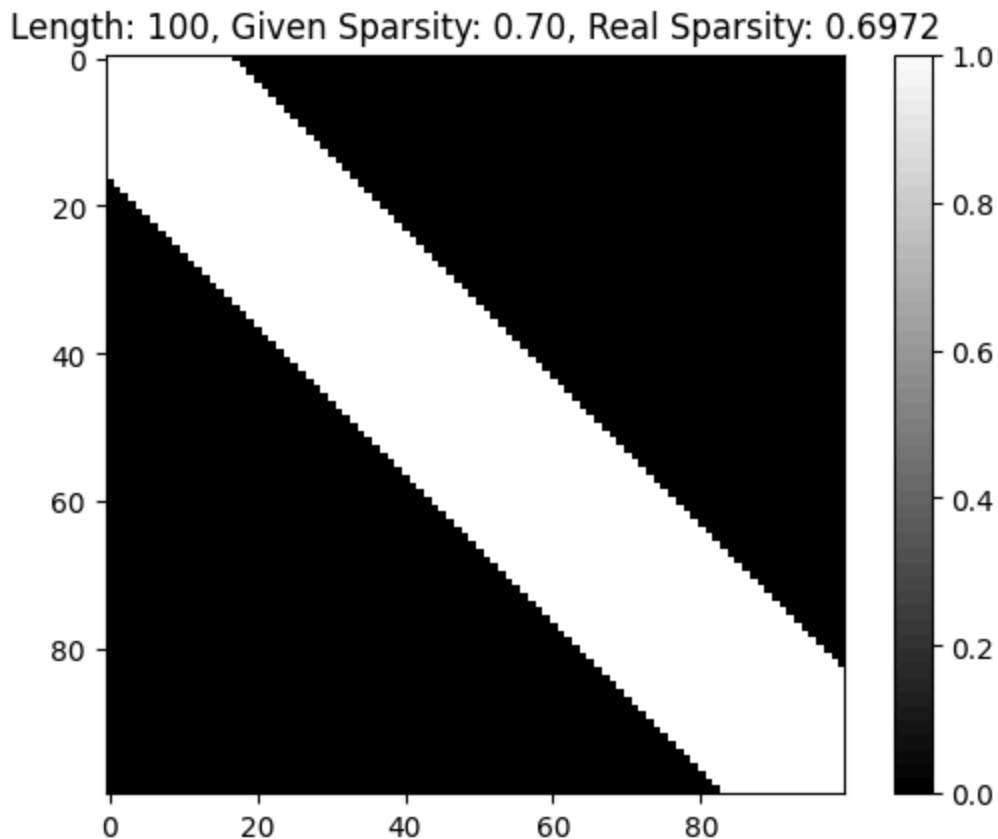


```
In [20]: def get_window_attention_mask_with_sparsity( length, sparsity):
# conditions : 0 <= sparsity <= 1
dw = best_diagonal_width_from_sparsity(length, sparsity)
return get_window_attention_mask( length=length, diagonal_width=dw)
```

```
In [21]: def generate_matrix_with_window_attention_mask_with_sparsity(length, sparsity):
# conditions : 0 <= diagonal_width <= 2*length - 1 (cover full matrix), diagonal_width is odd
matrix = np.ones((length, length))
mask = get_window_attention_mask_with_sparsity( length=length, sparsity=sparsity)
matrix[~mask] = 0
return matrix
```

```
In [22]: def test_generate_matrix_with_window_attention_mask_with_sparsity():
length = 100
sparsity = 0.7
matrix = generate_matrix_with_window_attention_mask_with_sparsity(length=length, sparsity=sparsity)
show_matrix_infos(matrix,length, sparsity)

test_generate_matrix_with_window_attention_mask_with_sparsity()
```



GLOBAL ATTENTION

Utils

```
In [23]: def global_attention_aera(length, global_attention_width):
w = global_attention_width
n = length
return (2 * w * n) - (w * w)

print(global_attention_aera(10,2))
```

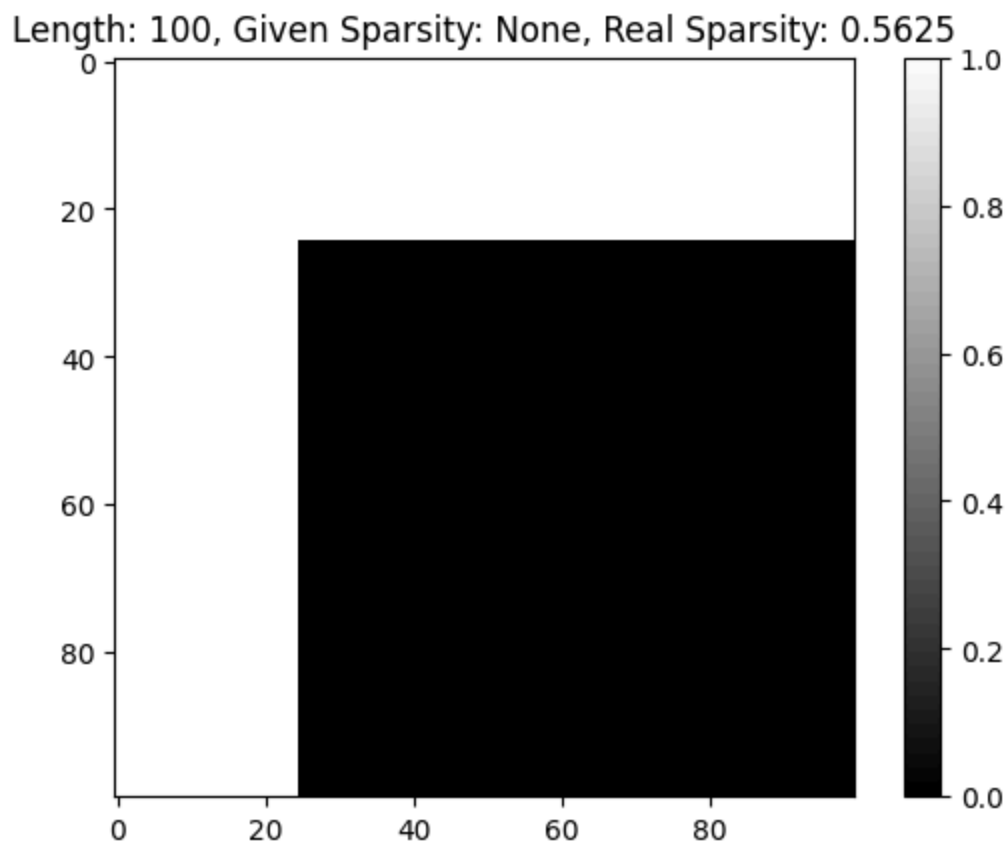
36

```
In [24]: def get_global_attention_mask( length, global_width):
mask = np.zeros(shape=(length,length), dtype=bool)
mask[:global_width,:] = True
mask[global_width : , : global_width] = True
return mask
```

```
In [25]: def generate_matrix_with_global_attention_mask(length, global_width):
matrix = np.ones((length, length))
mask = get_global_attention_mask( length=length, global_width=global_width)
matrix[~mask] = 0
return matrix
```

```
In [26]: def test_generate_matrix_with_global_attention_mask():
length = 100
global_width = 25
matrix = generate_matrix_with_global_attention_mask(length=length, global_width=global_width)
show_matrix_infos(matrix=matrix, length=length)

test_generate_matrix_with_global_attention_mask()
```

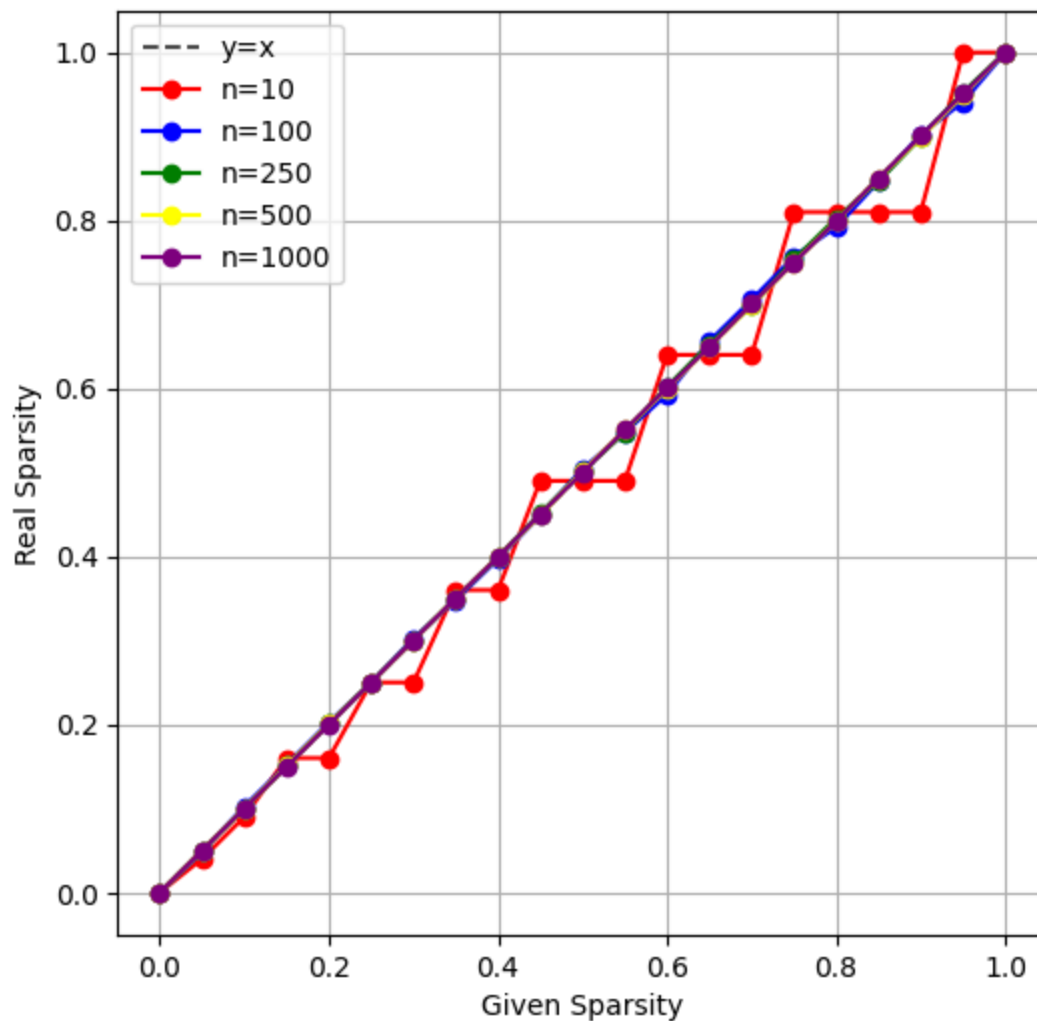


In [27]:

```
def best_global_width_from_sparsity(length, sparsity):
    n = length
    density = 1.0 - sparsity
    # ideal diagonal area
    ga = n * n * density
    # same as window mask but easier
    a = -1
    b = 2 * n
    c = -ga
    det = b * b - 4 * a * c
    x = (-b + math.sqrt(det)) / (2 * a)
    gw = round(x)
    if gw < 0: gw = 0
    elif gw > n * n: gw = n * n
    return gw
```

In [29]:

```
# test function hided for render
test_global_width_from_sparsity()
```

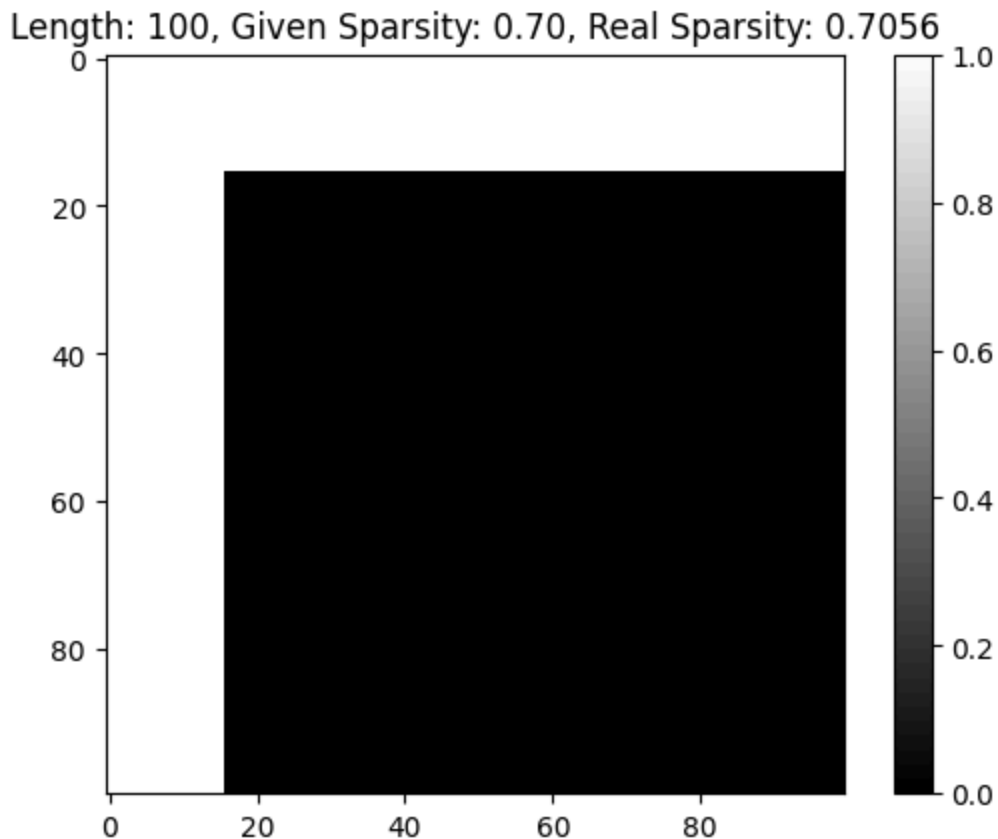



```
In [30]: def get_global_attention_mask_with_sparsity( length, sparsity):
# conditions : 0 <= sparsity <= 1
gw = best_global_width_from_sparsity(length, sparsity)
return get_global_attention_mask( length=length, global_width=gw)
```

```
In [31]: def generate_matrix_with_global_attention_mask_with_sparsity(length, sparsity):
matrix = np.ones((length, length))
mask = get_global_attention_mask_with_sparsity(length=length, sparsity=sparsity)
matrix[~mask] = 0
return matrix
```

```
In [32]: def test_generate_matrix_with_global_attention_mask_with_sparsity():
length = 100
sparsity = 0.7
matrix = generate_matrix_with_global_attention_mask_with_sparsity(length=length, sparsity=sparsity)
show_matrix_infos(matrix=matrix, length=length, given_sparsity= sparsity)

test_generate_matrix_with_global_attention_mask_with_sparsity()
```



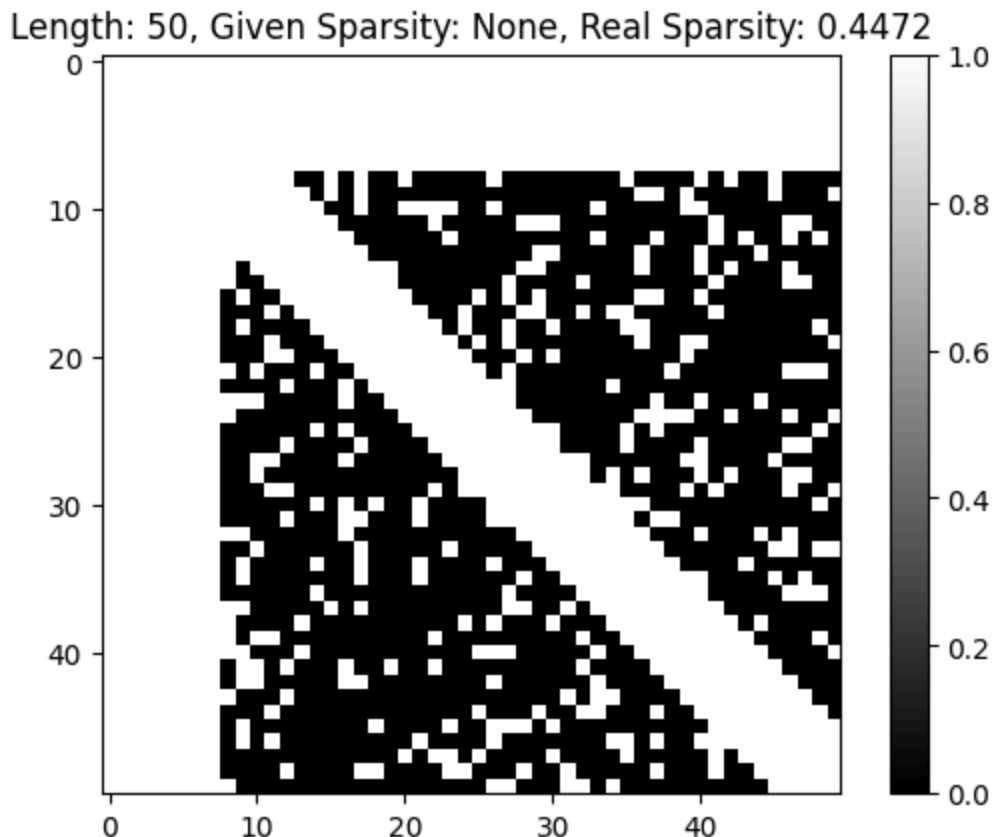
BIG BIRD (combination of all above)

```
In [33]: def get_big_bird_mask(length, nz_per_row, diagonal_width, global_width):
am = get_random_attention_mask( length=length, nz_per_row=nz_per_row)
wm = get_window_attention_mask( length=length, diagonal_width=diagonal_width)
gm = get_global_attention_mask( length=length, global_width=global_width)
total_mask = am | wm | gm
return total_mask
```

```
In [34]: def generate_big_bird(length, nz_per_row, diagonal_width, global_width ):
matrix = np.ones((length, length))
mask = get_big_bird_mask( length=length, nz_per_row=nz_per_row, diagonal_width=diagonal_width, global_width=global_wi
matrix[~mask] = 0
return matrix
```

```
In [35]: def test_generate_big_bird():
length = 50
nz_per_row = 10
diagonal_width = 8
global_width = 8
matrix = generate_big_bird(length=length,nz_per_row=nz_per_row, diagonal_width=diagonal_width, global_width=global_widt
show_matrix_infos(matrix=matrix, length=length)

test_generate_big_bird()
```

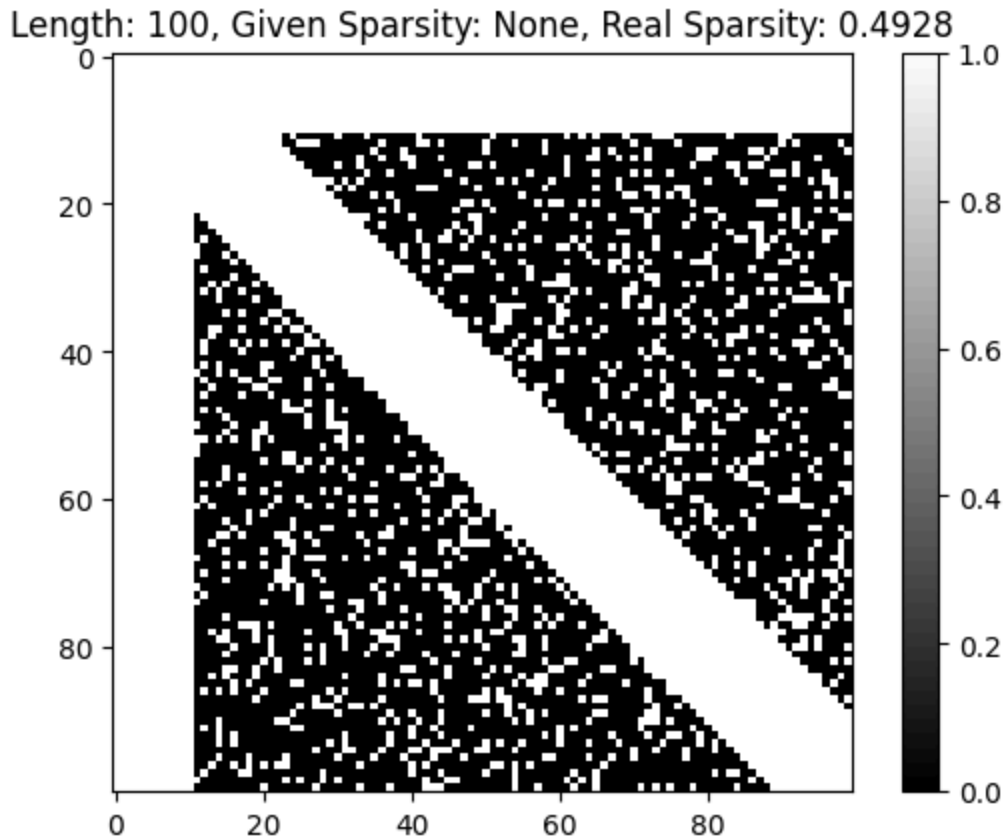


```
In [36]: def get_big_bird_mask_with_sparsity( length, random_sparsity, window_sparsity, global_sparsity):
am = get_random_attention_mask_with_sparsity( length= length, sparsity=random_sparsity)
wm = get_window_attention_mask_with_sparsity( length= length, sparsity=window_sparsity )
gm = get_global_attention_mask_with_sparsity(length=length, sparsity=global_sparsity)
total_mask = am | wm | gm
return total_mask
```

```
In [37]: def generate_big_bird_with_sparsity(length, random_sparsity, window_sparsity, global_sparsity):
matrix = np.ones((length, length))
mask = get_big_bird_mask_with_sparsity(length, random_sparsity, window_sparsity, global_sparsity)
matrix[~mask] = 0
return matrix
```

```
In [38]: def test_generate_big_bird_with_sparsity():
length = 100
random_sparsity = 0.8
window_sparsity = 0.8
global_sparsity = 0.8
matrix = generate_big_bird_with_sparsity(length=length, random_sparsity=random_sparsity, window_sparsity=window_sparsity)
show_matrix_infos(matrix=matrix, length= length)

test_generate_big_bird_with_sparsity()
```



```
In [39]: def adjust_total_sparsity(total_sparsity):
x = total_sparsity
# degree = 3
# a = 2.61815675
# b = -4.77052715
# c = 2.98999146
# d = 0.19945692
# res = a * ( x ** 3 ) + b * ( x ** 2 ) + c * x
# degree = 5
a = 24.08862473
b = -65.2963488
c = 64.48601296
d = -28.42365239
e = 5.98076684
f = 0.17082526
poly = a * x**5 + b * x**4 + c * x**3 + d * x**2 + e * x + f
res = min(max(poly, 0.0), 1.0)
return res
```

```
In [40]: def get_big_bird_mask_with_total_sparsity( length, total_sparsity, adjust):
if adjust :
total_sparsity = adjust_total_sparsity(total_sparsity)
random_sparsity = total_sparsity
window_sparsity = total_sparsity
global_sparsity = total_sparsity
total_mask = get_big_bird_mask_with_sparsity( length, random_sparsity, window_sparsity, global_sparsity)
return total_mask
```

```
In [41]: def generate_big_bird_with_total_sparsity(length,total_sparsity, adjust):
matrix = np.ones((length, length))
mask = get_big_bird_mask_with_total_sparsity(length, total_sparsity, adjust)
matrix[~mask] = 0
return matrix
```

```
In [42]: def find_approximation():
sparsity_values = [x / 100.0 for x in range(0, 101, 5)]
length = 2000

given_sparsity = []
real_sparsity = []

for sparsity in sparsity_values:
matrix = generate_big_bird_with_total_sparsity(length, sparsity, adjust=False)
real_sp = get_sparsity(matrix, length)
```

```

        given_sparsity.append(sparsity)
        real_sparsity.append(real_sp)

    sort_idx = np.argsort(real_sparsity)
    real_sorted = np.array(real_sparsity)[sort_idx]
    given_sorted = np.array(given_sparsity)[sort_idx]

    coeffs = np.polyfit(real_sorted, given_sorted, 5)

    print(f"a = {coeffs[0]}")
    print(f"b = {coeffs[1]}")
    print(f"c = {coeffs[2]}")
    print(f"d = {coeffs[3]}")
    print(f"e = {coeffs[4]}")
    print(f"f = {coeffs[5]}")
    print(f"poly = a * x**5 + b * x**4 + c * x**3 + d * x**2 + e * x + f")

    find_approximation()

```

a = 24.080837401839318

b = -65.27459395374147

c = 64.46369710621178

d = -28.413497112960528

e = 5.978905086888747

f = 0.17087493282757565

poly = a * x**5 + b * x**4 + c * x**3 + d * x**2 + e * x + f

```

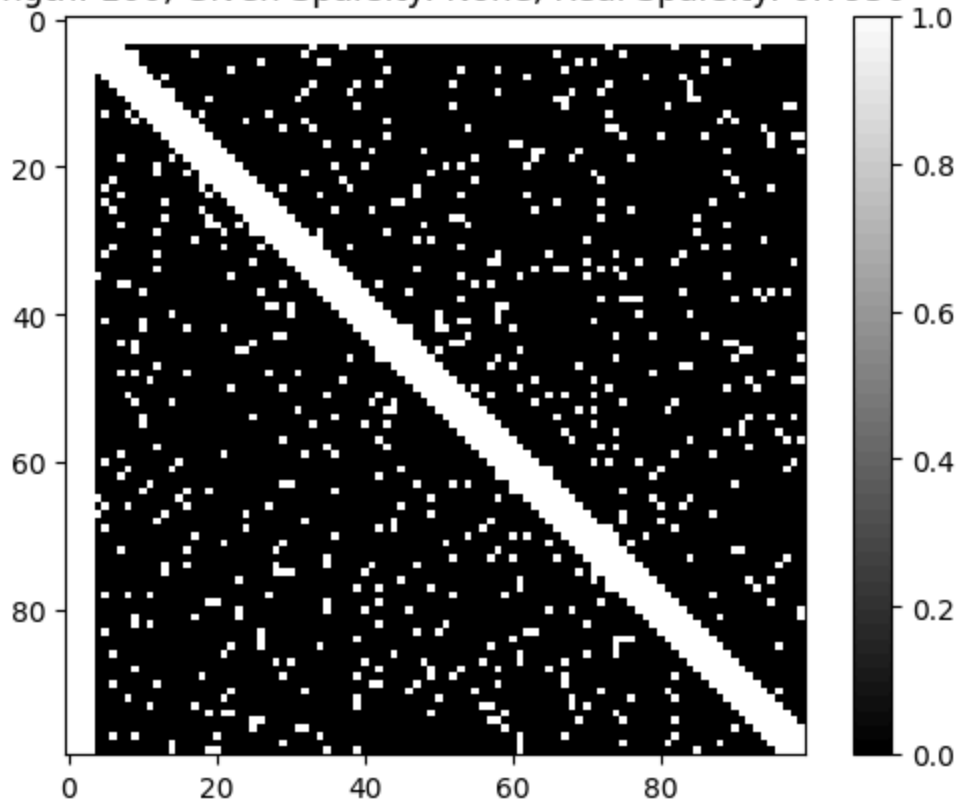
In [43]: def test_generate_big_bird_with_total_sparsity():
        length = 100
        total_sparsity = 0.8
        adjust = True
        matrix = generate_big_bird_with_total_sparsity(length, total_sparsity, adjust)
        print(f"Given sparsity = {total_sparsity}")
        show_matrix_infos(matrix, length)

        test_generate_big_bird_with_total_sparsity()

```

Given sparsity = 0.8

Length: 100, Given Sparsity: None, Real Sparsity: 0.7950



```

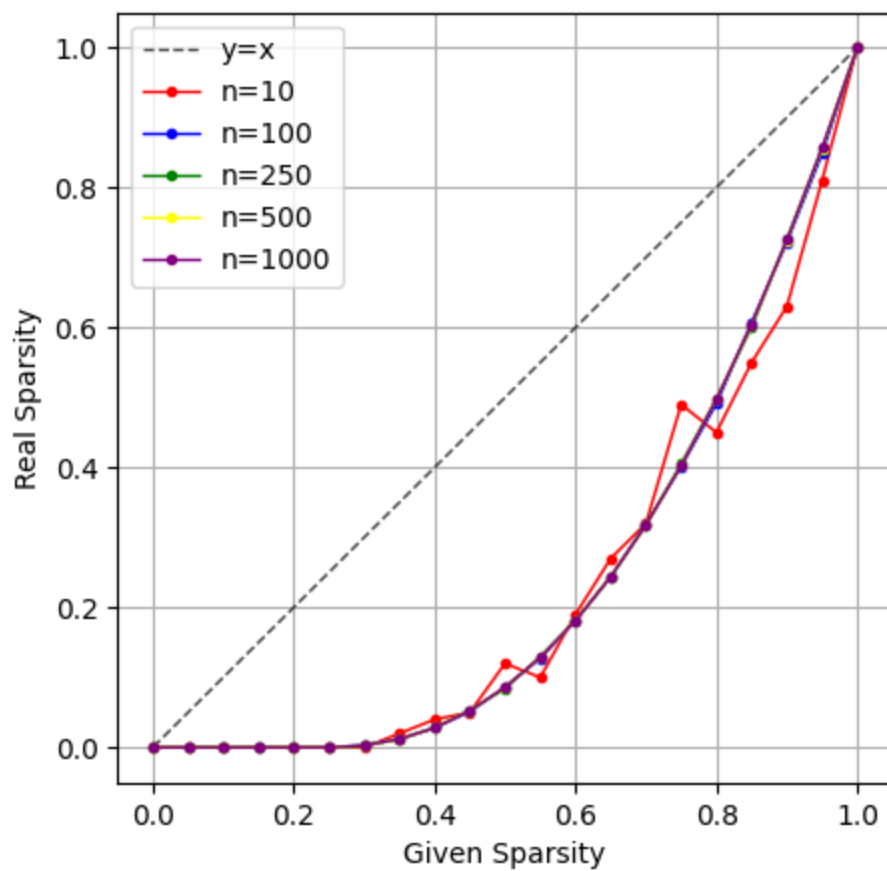
In [45]: # def test_adjust_total_sparsity(adjust):
        # ... (hidden for rendering)

        print("Without adjusting :")

```

```
test_adjust_total_sparsity(False)
print("When adjusting the given sparsity : ")
test_adjust_total_sparsity(True)
```

Without adjusting :



When adjusting the given sparsity :

